

SimData

Nick Miner

November 5, 2018

Question 1

1a

```
N <- 1000  
P <- 100
```

1b

```
Sigma = diag(sample(1:10, P, P))  
X <- mvrnorm(N, runif(P, -10, 10), Sigma)  
#Creating a 100x100 matrix for Sigma with a random, repeat sampling of 1 through 10 running through  
#the matrix diagonal. Setting X as a multivariate normal with 100 obs, a uniform mean of 0, and Sigma  
#as the covariance matrix for the data.
```

1c

```
p <- rbinom(P, 1, 0.1)  
#vector p from Bernoulli distribution.
```

1d

```
beta = p * rnorm(P, 5, sqrt(5)) + (1 - p) * rnorm(P, 0, sqrt(0.1))  
epsilon <- rnorm(N, 0, sqrt(10))  
y <- X %*% beta + epsilon  
#Creating beta to draw from different normal distributions based on if p is 1 or 0. Setting epsilon  
#as a normal distribution with mean 0 and sd of sqrt(10). Creating y as a function of X times beta  
#plus error term epsilon.
```

1e

```
dataset1 <- as.data.frame(cbind(y,X))  
names(dataset1) <- c("y", rep(paste("X", 1:ncol(X), sep = "")))  
train.index <- sample(1:1000, size = 800, replace = FALSE)  
test.index <- setdiff(1:1000, train.index)  
training.set <- dataset1[train.index,]  
test.set <- dataset1[test.index,]  
#Creating a dataset from x and Y, then setting the colnames for the dataset. Creating training and  
#test sets from a random sampling of the data without overlap.
```

Question 2

2a

```
q2.lm <- lm(y ~ X)
coef(summary(q2.lm))[1:20,]
```

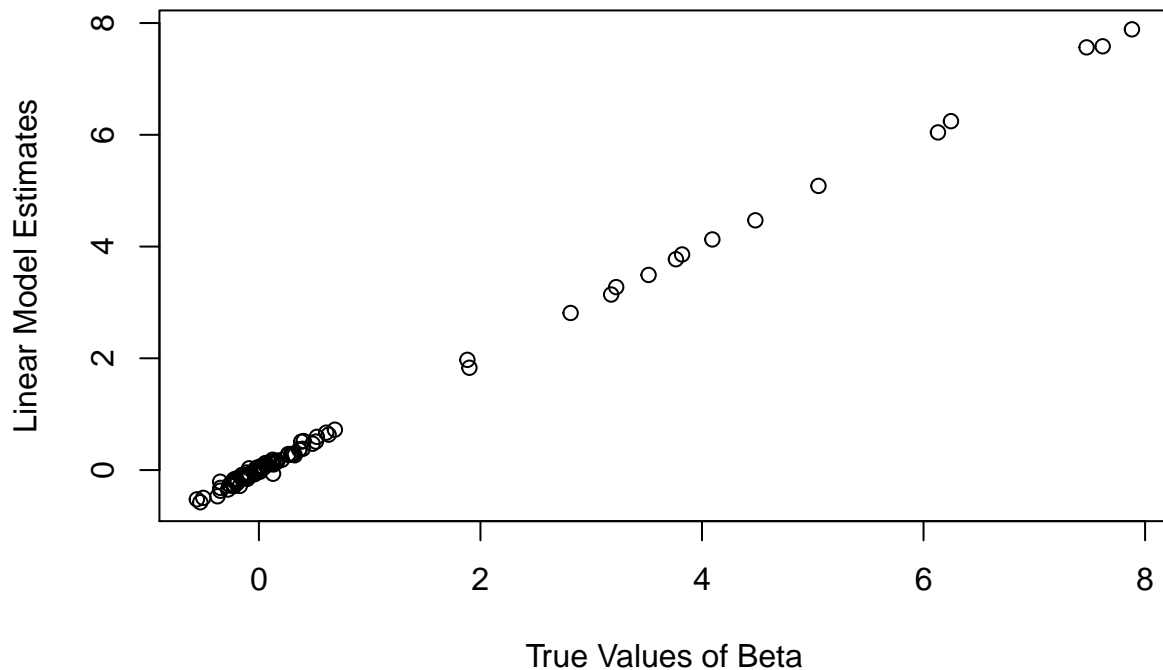
##	Estimate	Std. Error	t value	Pr(> t)
## (Intercept)	0.050508616	3.34853995	0.01508377	9.879687e-01
## X1	0.112978844	0.03829196	2.95045898	3.255360e-03
## X2	-0.278295706	0.03492325	-7.96877992	4.838691e-15
## X3	-0.240938027	0.03870069	-6.22567818	7.348743e-10
## X4	-0.156600438	0.03539987	-4.42375719	1.088805e-05
## X5	0.121391724	0.04971833	2.44158894	1.481466e-02
## X6	-0.286864608	0.07535648	-3.80676773	1.503404e-04
## X7	0.031884159	0.05450956	0.58492778	5.587430e-01
## X8	0.722684497	0.04271400	16.91914912	5.862950e-56
## X9	-0.008950315	0.03836953	-0.23326622	8.156078e-01
## X10	-0.070981390	0.03591945	-1.97612704	4.844575e-02
## X11	-0.348579149	0.10606934	-3.28633269	1.054317e-03
## X12	-0.193412405	0.07535315	-2.56674615	1.042662e-02
## X13	0.163947684	0.03701727	4.42895138	1.063460e-05
## X14	-0.065662045	0.10425569	-0.62981736	5.289741e-01
## X15	0.044494810	0.05442269	0.81757825	4.138147e-01
## X16	1.831336774	0.04816250	38.02412041	2.394697e-189
## X17	0.124612878	0.05335538	2.33552621	1.973507e-02
## X18	7.583933312	0.04788321	158.38398492	0.000000e+00
## X19	-0.575859709	0.07589338	-7.58774599	8.127662e-14

```
#Fitting linear model and examining the first 20 coefficients.
```

```
#Plotting our data generated beta values against linear model  
#estimates. Plotting the data shows a perfect match of the linear  
#estimates to the true values of beta.
```

```
plot(beta, coef(q2.lm)[-1], main = "Magnitude of True Beta Parameters vs Linear Estimates",  
      xlab = "True Values of Beta", ylab = "Linear Model Estimates")
```

Magnitude of True Beta Parameters vs Linear Estimates



2b

#To test the predictive accuracy of the model, I fit the linear model to my previously created training dataset. Then I found the mse of that model to be 8.93. I then fit the training model to the previously created test set to create test predictions. I found the mse of the test predictions to be 12.21. The training and test mse are reasonably close and reasonably low. Still, a difference of over 3 seems to indicate the model might not be the best fit.

```
training.model <- lm(y~., data = training.set)
training.mse <- mean(((training.set$y - training.model$fitted.values)^2))
training.mse
```

```
## [1] 8.934155
```

#setting a linear model to the previously created training data. Then setting and viewing the mean squared error of the test set as a manner for assessing model predictive accuracy.

```
test.preds <- predict(training.model, newdata = test.set)
test.mse <- mean(((test.set$y - test.preds)^2))
test.mse
```

```
## [1] 12.2089
```

#Predicting values for the test set using the model fit on the training set. Then finding the mean squared error of the test predictions.

Question 3

3a

*#Penalized regression is used in cases where the number of predictors $p >$ the number of obs n .
#Penalized regression introduces some form of penalty based on the beta coefficients, multiplied
#by a regularization parameter λ that controls the penalty vs model fit. This penalty shrinks
#beta coefficients towards zero, thus constraining the abundance of parameters and allowing for a
#model to be fit to the data.*

3b

*#In a LASSO regression, the regression penalty is the sum of the absolute value of β_j . Multiplied
#by regulator λ , this regression approaches the target point by setting the mean of the
#distribution at 0 and narrowing the parameters diagonally. LASSO regressions don't have as many nice
#features as ridge regressions, but shrinks many parameters to exactly 0, which can lead to cleaner
#results.*

*#For a ridge regression the penalty is the sum of squared β_j multiplied by regulator λ .
#Ridge regression sets the mean of the distribution at 1 and approaches the target point circularly.
#Ridge regressions are useful for data with multicollinearity, optimize more easily, and are able to
#be smoothed for stable results.*

*#Lastly, an Elastic Net regression uses a combination of both Ridge and LASSO. Both penalties are
#multiplied by the same λ , but then also multiplied by an α value between 0 and 1. The sum
#of the penalty values selected for by one value of α are added to the sum of the values for the
#opposing regression multiplied by $1 - \alpha$ to ensure independent draws. Elastic Net regressions
#linearly combine the Ridge and LASSO regressions for the most efficient way to constrain parameters.
#The Ridge Regression shrinks most parameters close to zero while the LASSO Regression shrinks most
#parameters to zero. The Elastic Net draws from each based on probability selector α , and then
#sums the values produced by both regressions selected by α . Because two penalties are applied,
#there can be "double-shrinkage"; ie the model may have shrank the selection area too quickly and
#missed the target value. To combat this the Elastic Net Regression uses a quadratic penalty. The
#Elastic Net regression combines the best features of LASSO and Ridge regressions while also
#accounting for new penalties more effectively.*

3c

*#My own research was to examine the combat effectiveness of Marines in Iraq from 2001-2009. However,
#I would have to account for numerous parameters, such as time of year, weather, number of combat
#veterans per squad, number of new recruits per squad, gear variations, etc. Since my data would be
#coming from the DOD, there is a strong chance my parameters would far exceed my observations. A
#linear model would be extremely ineffective, if not entirely useless. A penalized regression model,
#however, may help to control the parameters so effective results can be observed.*

Question 4

4a

```
mod.lasso = train(y~., method = "glmnet", tuneGrid = expand.grid(alpha = 0, lambda = seq(0, 50, .5)),
                  data = dataset1, preProcess = c("center"), trControl = trainControl(method = "cv",
                                                                                      number = 2,
                                                                                      search = "grid"))

#fitting LASSO model with broad search spectrum
mod.lasso$bestTune

##      alpha lambda
## 5         0      2

#best fit alpha and lambda for search area

mod.lasso2 = train(y~., method = "glmnet", tuneGrid = expand.grid(alpha = 0, lambda = seq(1, 3, .1)),
                   data = dataset1, preProcess = c("center"), trControl = trainControl(method = "cv",
                                                                                       number = 2,
                                                                                       search = "grid"))


#re-running LASSO with lambda parameters reset to above and below lambda value found from first run.
mod.lasso2$bestTune

##      alpha lambda
## 12         0     2.1

#showing new alpha and lambda

#Creating predictions for y based off of second LASSO run, then plotting those y values vs the data
#generated y values. The plot shows the predicted y values and actual y values to be an extremely
#close fit, with only a few outliers. The mse for the model is calculated as 14.40.
lasso.y.preds <- predict(mod.lasso2)
plot(dataset1$y, lasso.y.preds, main = "LASSO Predicted Y Values vs True Y Values",
      xlab = "True Y Values", ylab = "LASSO Predicted Y Values")
```

LASSO Predicted Y Values vs True Y Values



```
lasso.mse <- mean(((dataset1$y - lasso.y.preds)^2))
lasso.mse
```

```
## [1] 14.39556
```

4b

```
mod.ridge = train(y~., method = "glmnet", tuneGrid = expand.grid(alpha = 1, lambda = seq(0, 50, .5)),
                  data = dataset1, preprocess = c("center"), trControl = trainControl(method = "cv",
                                                                                       number = 2,
                                                                                       search = "grid"))
```

```
## Warning in nominalTrainWorkflow(x = x, y = y, wts = weights, info =
## trainInfo, : There were missing values in resampled performance measures.
```

```
#setting broad ridge
mod.ridge$bestTune
```

```
##   alpha lambda
## 1      1      0
```

```
#showing alpha and lambda
```

```
mod.ridge2 = train(y~., method = "glmnet", tuneGrid = expand.grid(alpha = 1, lambda = seq(-1, 1, .1)),
                   data = dataset1, preprocess = c("center"), trControl = trainControl(method = "cv",
                                                                                       number = 2,
                                                                                       search = "grid"))
```


```
#second ridge run with new parameters
mod.ridge2$bestTune
```

```
##      alpha lambda
## 11      1      0
```

```
#new alpha and lambda
```

```
#Creating y predictions based on the second ridge model and then plotting those predictions vs the
#data generated y values. The Ridge regression shows a model fit as least as tight as the LASSO, and
#actually has a lower mse: 9.34 vs 14.40.
```

```
ridge.y.preds <- predict(mod.ridge2)
plot(dataset1$y, ridge.y.preds, main = "Ridge Predicted Y Values vs True Y Values",
      xlab = "True Y Values", ylab = "Ridge Predicted Y Values")
```



```
ridge.mse <- mean(((dataset1$y - ridge.y.preds)^2))
ridge.mse
```

```
## [1] 9.340122
```

4c

```
mod.elastic.net = train(y~., method = "glmnet", tuneGrid = expand.grid(alpha = .5,
                                                                    lambda = seq(0, 50, 1)),
                        data = dataset1, preprocess = c("center"), trControl = trainControl(method = "cv",
```

```

number = 2,
search = "grid"))

## Warning in nominalTrainWorkflow(x = x, y = y, wts = weights, info =
## trainInfo, : There were missing values in resampled performance measures.
#Fitting elastic net model. Alpha is set as .5 for even selection from Ridge and LASSO penalties.
mod.elastic.net$bestTune

##   alpha lambda
## 1    0.5      0
#finding best fit alpha and lambda.

mod.elastic.net2 = train(y~., method = "glmnet", tuneGrid = expand.grid(alpha = .5,
                                                                    lambda = seq(0, 2, .1)),
                        data = dataset1, preProcess = c("center"), trControl = trainControl(method = "cv",
                                                                    number = 2,
                                                                    search = "grid"))


#running elastic net second time.
mod.elastic.net2$bestTune

##   alpha lambda
## 1    0.5      0
#new alpha and lambda

#Generating y predictions based off of the second elastic net run and plotting those values against the
#data generated y values. This plot shows the model fits even more tightly than the LASSO and Ridge
#regressions. The model also has one of the lowest mse's at 9.35.
elastic.net.y.preds <- predict(mod.elastic.net2)
plot(dataset1$y, elastic.net.y.preds, main = "Elastic Net Predicted Y Values vs True Y Values",
      xlab = "True Y Values", ylab = "Elastic Net Predicted Y Values")

```


Elastic Net Predicted Y Values vs True Y Values



```
elastic.net.mse <- mean(((dataset1$y - elastic.net.y.preds)^2))
elastic.net.mse
```

```
## [1] 9.350061
```

4d

#All 3 of the models performed far superiorly to the linear model. Because the models are designed #to penalize additional betas produced by parameters when $p > n$, the models constrain the parameters #to allow for regression in a way that linear regressions do not. Essentially penalized regressions #constrain excess parameters to allow for linear regression through the penalties.

#The best model appears to be the Ridge regression, in terms of predicted y values to DGP y values. #The LASSO model performs the worst, and the Elastic Net regression performs only slightly worse than #the Ridge regression. Because the Elastic Net regression utilizes the LASSO penalty in part, and the #LASSO model performs poorly on collinear data, it appears that this data is somewhat collinear.

Question 5

5a

#The linear model should not work in this case. Since the number of predictors p is larger than the #number of observations n there may be an infinite number of modeling options for the model. This

#also means there is no longer a unique least-squares coefficient estimate since variance for data #points is infinite. Below is a repeat of the data generating process from question 2, but with p = #1500.

```
N <- 1000
P <- 1500
Sigma = diag(sample(1:10, P, P))
X <- mvrnorm(N, runif(P, -10, 10), Sigma)
p <- rbinom(P, 1, 0.1)
beta = p * rnorm(P, 5, sqrt(5)) + (1 - p) * rnorm(P, 0, sqrt(0.1))
epsilon <- rnorm(N, 0, sqrt(10))
y <- X %*% beta + epsilon

dataset2 <- as.data.frame(cbind(y,X))
names(dataset2) <- c("y", rep(paste("X", 1:ncol(X), sep = "")))
train.index2 <- sample(1:1000, size = 800, replace = FALSE)
test.index2 <- setdiff(1:1000, train.index2)
training.set2 <- dataset2[train.index2,]
test.set2 <- dataset2[test.index2,]
```

5b

#The summary of the linear model reports that all 1000 residuals are 0 with 0 degrees of freedom. While #the variable coefficients have values, there is no standard error, t-values or probabilities reported #for any variables. There are no residuals, degrees of freedom or anything other than coefficient values. #because the number of parameters forced the data into higher dimensions, resulting in the data being #too sparse to attribute probabilities correctly to the model.

```
q5.lm <- lm(y~X)
coef(summary(q5.lm))[1:20,]
```

##	Estimate	Std. Error	t value	Pr(> t)
## (Intercept)	15271.718150	NaN	NaN	NaN
## X1	-113.434175	NaN	NaN	NaN
## X2	9.804141	NaN	NaN	NaN
## X3	143.917326	NaN	NaN	NaN
## X4	-11.386549	NaN	NaN	NaN
## X5	-244.845930	NaN	NaN	NaN
## X6	75.448846	NaN	NaN	NaN
## X7	251.931083	NaN	NaN	NaN
## X8	11.260908	NaN	NaN	NaN
## X9	173.432319	NaN	NaN	NaN
## X10	44.952362	NaN	NaN	NaN
## X11	-5.501978	NaN	NaN	NaN
## X12	19.846197	NaN	NaN	NaN
## X13	-49.382156	NaN	NaN	NaN
## X14	-411.450426	NaN	NaN	NaN
## X15	70.578495	NaN	NaN	NaN
## X16	-157.511800	NaN	NaN	NaN
## X17	-18.856145	NaN	NaN	NaN
## X18	28.924985	NaN	NaN	NaN
## X19	76.399507	NaN	NaN	NaN

5c

5c-LASSO

```
q5.cv.lasso <- cv.glmnet(x = X, y = y, alpha = 0)
#using cross-validation to find the optimal starting lambda value for LASSO.

q5.cv.lasso$lambda.min

## [1] 391.6764

q5.cv.lasso$lambda.1se

## [1] 471.7756
#Finding the minimum lambda value and 1 se value.

q5.mod.lasso = train(y~., method = "glmnet", tuneGrid = expand.grid(alpha = 0, lambda = seq(0, 6, .1)),
                    data = dataset2, preprocess = c("center"), trControl = trainControl(method = "cv",
                                                                                          number = 2,
                                                                                          search = "grid"))

#Setting parameters for lambda as the value of 1 se above and below the minimum lambda value, rounded
#to the nearest tenth.
q5.mod.lasso$bestTune

##      alpha lambda
## 61      0      6
#new alpha and lambda

q5.mod.lasso2 = train(y~., method = "glmnet", tuneGrid = expand.grid(alpha = 0,
                                                                    lambda = seq(3, 9, .001)),
                    data = dataset2, preprocess = c("center"), trControl = trainControl(method = "cv",
                                                                                          number = 2,
                                                                                          search = "grid"))


#setting new lambda and search as dictated by 1 se.
q5.mod.lasso2$bestTune

##      alpha lambda
## 6001      0      9
#new alpha and lambda.

#Predicting y values determined by the second LASSO run, then plotting those values against the data
#generated y values. The plot shows the model does fit, but not nearly as neatly as in the previous
#question. In addition, the mse is 7129.41, which is incredibly high compared to the previous models'
#mse. The LASSO model's poor fit to the data might indicate that the data is multicollinear, and
#therefore the LASSO is ignoring the data clustering when it draws.

q5.lasso.y.preds <- predict(q5.mod.lasso2)
plot(dataset2$y, q5.lasso.y.preds, main = "LASSO Predicted Y Values vs True Y Values",
     xlab = "True Y Values", ylab = "LASSO Predicted Y Values")
```

LASSO Predicted Y Values vs True Y Values



```
q5.lasso.mse <- mean(((dataset2$y - q5.lasso.y.preds)^2))
q5.lasso.mse
```

```
## [1] 7514.088
```

5c-Ridge

```
q5.cv.ridge <- cv.glmnet(x = X, y = y, alpha = 1)
#using cross-validation to determine optimal starting lambda value
```

```
q5.cv.ridge$lambda.min
```

```
## [1] 0.4103264
```

```
q5.cv.ridge$lambda.1se
```

```
## [1] 0.4942395
```

```
#finding the minimum lambda value and the value of 1 se for the model.
```

```
q5.mod.ridge = train(y~., method = "glmnet", tuneGrid = expand.grid(alpha = 1,
                                                                    lambda = seq(-0.05, 0.2, .001)),
                    data = dataset2, preprocess = c("center"), trControl = trainControl(method = "cv",
                                                                                       number = 2,
                                                                                       search = "grid"))
```

```
#Ridge regression with starting lambda search.
```

```

q5.mod.ridge$bestTune

##      alpha lambda
## 251      1    0.2
#new alpha and lambda

q5.mod.ridge2 = train(y~., method = "glmnet", tuneGrid = expand.grid(alpha = 1,
                                                                    lambda = seq(0.1, 0.3, 0.01)),
                    data = dataset2, preprocess = c("center"), trControl = trainControl(method = "cv",
                                                                                          number = 2,
                                                                                          search = "grid"))


#second ridge run with new lambda parameters.
q5.mod.ridge2$bestTune

##      alpha lambda
## 21      1    0.3
#new alpha and lambda.

#This plot shows the ridge model fits the data much more neatly than the LASSO model. Also, the mse
#of the model is 80.23, which is significantly smaller than the LASSO's mse.

q5.ridge.y.preds <- predict(q5.mod.ridge2)
plot(dataset2$y, q5.ridge.y.preds, main = "Ridge Predicted Y Values vs True Y Values",
     xlab = "True Y Values", ylab = "Ridge Predicted Y Values")

```



```
q5.ridge.mse <- mean(((dataset2$y - q5.ridge.y.preds)^2))
q5.ridge.mse
```

```
## [1] 100.4291
```

5c-ENet

```
q5.cv.enet <- cv.glmnet(x = X, y = y, alpha = 0.5)
#Using cross-validation to discover starting lambda
```

```
q5.cv.enet$lambda.min
```

```
## [1] 0.7477482
```

```
q5.cv.enet$lambda.1se
```

```
## [1] 0.8206527
```

```
#minimum lambda value and the value for 1 se for the model.
```

```
q5.mod.enet = train(y~., method = "glmnet", tuneGrid = expand.grid(alpha = 0.5,
                                                                    lambda = seq(-0.08, 0.3, 0.01)),
                    data = dataset2, preProcess = c("center"), trControl = trainControl(method = "cv",
                                                                                          number = 2,
                                                                                          search = "grid"))
```

```
#Fitting Elastic Net model to data with search parameters around starting lambda.
```

```
q5.mod.enet$bestTune
```

```
##      alpha lambda
## 39    0.5    0.3
```

```
#new alpha and lambda
```

```
q5.mod.enet2 = train(y~., method = "glmnet", tuneGrid = expand.grid(alpha = .5,
                                                                    lambda = seq(0.2, 0.4, .001)),
                    data = dataset2, preProcess = c("center"), trControl = trainControl(method = "cv",
                                                                                          number = 2,
                                                                                          search = "grid"))
```

```
#New Elastic Net model with lambda parameters set around new value.
```

```
q5.mod.enet2$bestTune
```

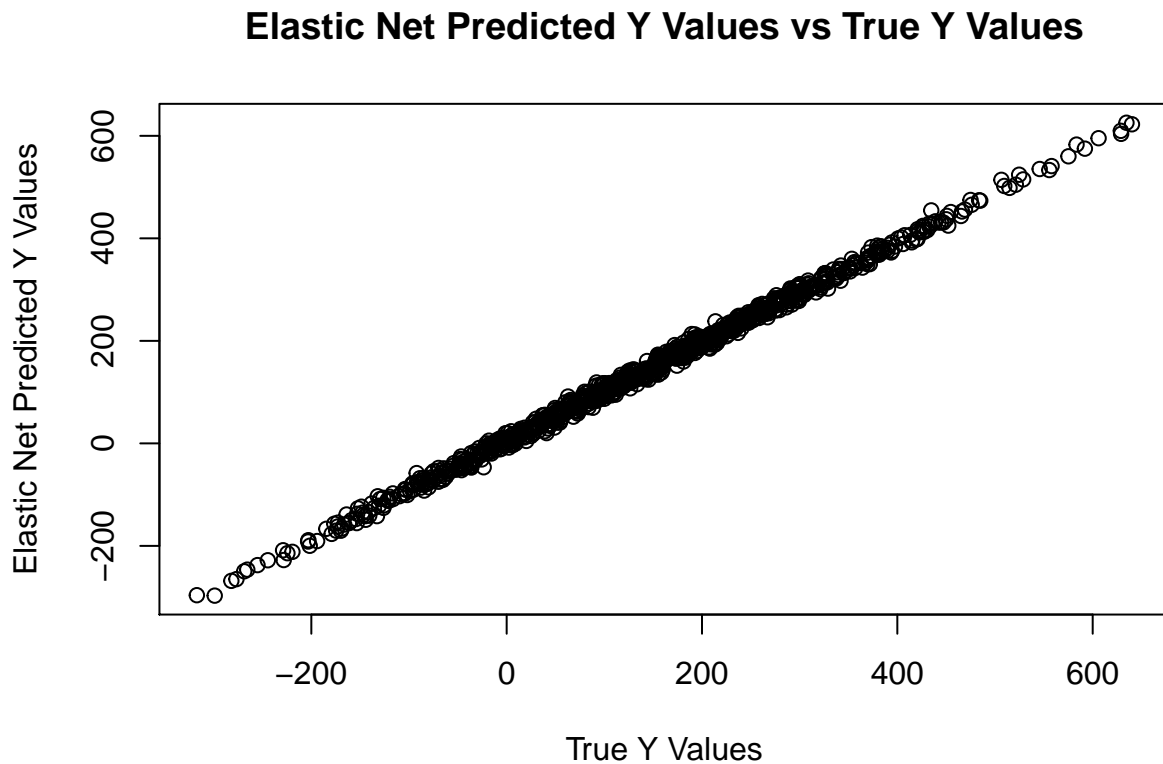
```
##      alpha lambda
## 201    0.5    0.4
```

```
#new alpha and lambda
```

```
#The plot shows this model also fits the data much more neatly than the LASSO model, although not quite
#as neatly as the Ridge model. Since the Elastic Net model relies on penalties from the other 2 models,
#the penalties from the poorly-fit LASSO model may be causing some differences in the beta coefficients
#causing the Elastic Net model to be a slightly poorer fit than just the Ridge model. The mse for the
#Elastic Net is 83.11, which is only slightly higher than the Ridge model's. Elastic Net models also
#work better with a larger number of observations, so 1000 obs may not have been enough.
```

```
q5.enet.y.preds <- predict(q5.mod.enet2)
```

```
plot(dataset2$y, q5.enet.y.preds, main = "Elastic Net Predicted Y Values vs True Y Values",
      xlab = "True Y Values", ylab = "Elastic Net Predicted Y Values")
```



```
q5.enet.mse <- mean(((dataset2$y - q5.enet.y.preds)^2))
q5.enet.mse
```

```
## [1] 104.1902
```

5d


5d-LASSO

#The predicted y values vs data generated y values were examined in question 5c, so in question 5d I will examine the true beta values vs the model estimates.

#The beta estimates from the LASSO model plotted against the DGP beta values shows the model is an excellent fit for the data. There is only one extreme outlier, while the rest of the points lie along the axis. The plot does show many points clustered around (0,0), indicating possible multicollinearity in the data.

```
plot(sample(beta, size = length(coef(q5.mod.lasso2$finalModel, q5.mod.lasso2$bestTune$lambda)),
      replace = TRUE), coef(q5.mod.lasso2$finalModel, q5.mod.lasso2$bestTune$lambda),
      main = "Magnitude of True Beta Parameters vs Linear Estimates",
      xlab = "True Values of Beta", ylab = "Linear Model Estimates")
```

Magnitude of True Beta Parameters vs Linear Estimates




5d-Ridge

#The ridge plot is very similar to the LASSO plot. There is only one extreme outlier and much of the #data appears clustered around (0,0).

```
plot(sample(beta, size = length(coef(q5.mod.ridge2$finalModel, q5.mod.ridge2$bestTune$lambda)),
         replace = TRUE), coef(q5.mod.ridge2$finalModel, q5.mod.ridge2$bestTune$lambda),
     main = "Magnitude of True Beta Parameters vs Linear Estimates",
     xlab = "True Values of Beta", ylab = "Linear Model Estimates")
```


Magnitude of True Beta Parameters vs Linear Estimates

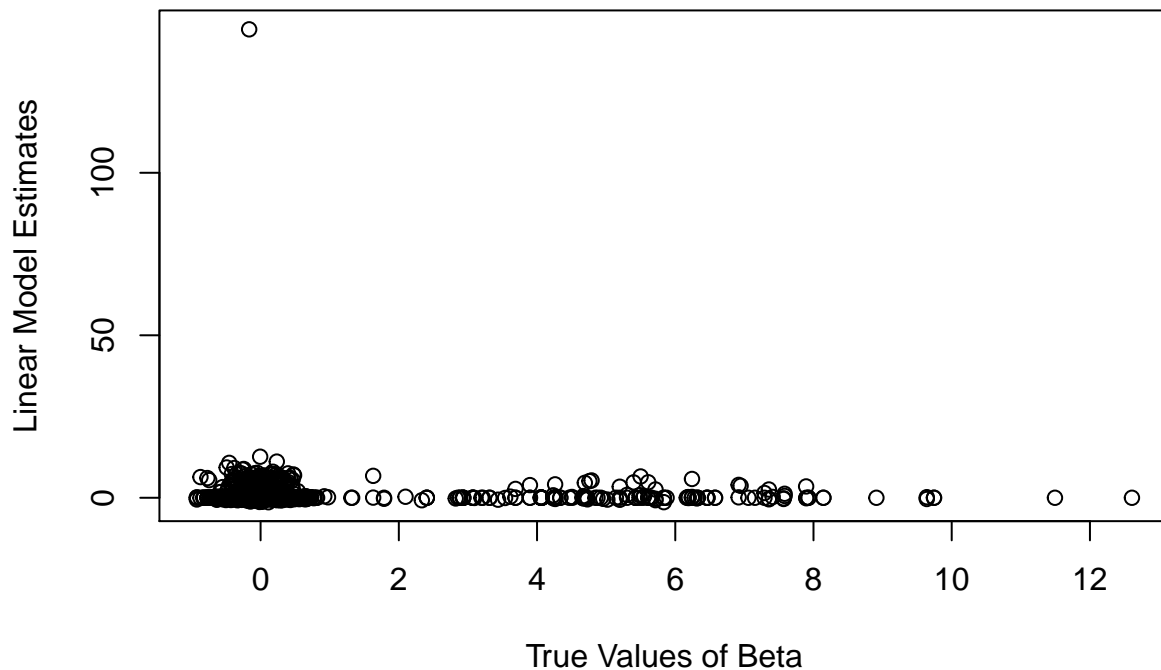


5d-ENet

#The Elastic Net follows the same pattern as the other plots, with one extreme outlier and much of the data centered on (0,0).

```
plot(sample(beta, size = length(coef(q5.mod.enet2$finalModel, q5.mod.enet2$bestTune$lambda)),
         replace = TRUE), coef(q5.mod.enet2$finalModel, q5.mod.enet2$bestTune$lambda),
     main = "Magnitude of True Beta Parameters vs Linear Estimates",
     xlab = "True Values of Beta", ylab = "Linear Model Estimates")
```

Magnitude of True Beta Parameters vs Linear Estimates



5e

#For this procedure, I had a dataset consisting of 1000 observations n and 1500 predictors P . The data generating process was comprised from $Y = X \cdot \text{Beta} + \text{Epsilon}$. Y is the resulting dataframe of multivariate matrix X (drawn from a uniform distribution and multiplied diagonally against matrix Sigma multiplied by Beta (drawn from 2 different normal distributions based on propensity p) with an added error term (drawn from a normal distribution).

#I divided the resulting dataframe Y into training and test sets. I applied a linear model to the data, but the linear model reported no results. Because the number of predictors P outnumbers the number of obs n , there may be an infinite number of modeling options for the model. This also means there is no longer a unique least-squares coefficient estimate since variance for data points is infinite.

#Afterwards I applied LASSO, Ridge, and Elastic Net regression models to the data. This was necessary because the aforementioned models apply a penalized regression to the data. This penalized regression helps to constrain additional parameters to allow for linear regression to be accomplished. I began with the LASSO model, which produced predicted y values close to the DGP y values. However, the SE of the model was incredibly high, and the y value and beta value plotting indicated another model may perform better. The Ridge model fit the data much more tightly in the y and beta values plots, and had a significantly smaller SE. Finally, the Elastic Net model performed better than the LASSO model, but not quite as well as the Ridge regression. The LASSO regression deals poorly with collinearity, ignoring data clustering on its draws. Since the LASSO regression and Elastic Net regression (which uses the LASSO penalty in part) both performed more poorly on the data than the Ridge regression, it seems the data is in some part collinear.

*#Improving on this model would include finding a penalized regression technique that performs well
#on collinear data. A regression only using that model could be performed, and then the Elastic Net
#regression could be re-used with that regression method's penalty replacing the LASSO penalty.
#Then another comparison could be drawn between the models to evaluate accuracy.*